

# Implementing GANs

Duke MLSS

Alex Lew

# Probability Distributions

$$x \sim p(x)$$

$x$  “is drawn from”  $p(x)$

$x$  “is a sample from”  $p(x)$

$p(x)$  a distribution assigning 50% probability to H, 50% to T

$x$  a random variable, equal to H or T



# Probability Distributions

$$x \sim p(x)$$

$x$  “is drawn from”  $p(x)$

$x$  “is a sample from”  $p(x)$



$p(x)$  a distribution assigning some small probability to each possible setting of pixel values in a 24-bit-color 100x100 image.

$x$  a random variable, equal to some 100x100 image.

# Generative modeling

Approximate  $p(x)$  that generated the data

Do we want to...

- evaluate relative likelihoods of different data points?
- compute conditional distributions?
- do inference on latent variables?
- produce new samples (perhaps conditioned on something)?

Note: we have no “labeled”  $p(x)$  values,  
just lots of  $x$

# Generative modeling

For today:  
**Generate New Samples**

(create “realistic” images, sounds, etc.)

# Sampling

How do we generate a sample from a distribution?

Method 1: Go out and flip a coin.  
(slow, expensive)

$p(x)$  a distribution assigning 50% probability to H, 50% to T  
 $x$  a random variable, equal to H or T



# Sampling

How do we generate a sample from a distribution?

Method 2: *Simulate* a sample.

`np.random.uniform(0,1)`



Source of (pseudo-)randomness

```
def G(z):  
    if z > 0.5:  
        return "Heads"  
    else:  
        return "Tails"
```



Deterministic mapping



Sample from  
 $p(x)$

# Sampling

But sometimes, we don't know exactly what  $G(z)$  should be.





# Sampling

But sometimes, we don't know exactly what  $G(z)$  should be.

We can estimate *parameters* from *data*.

`np.random.uniform(0,1)`



Source of (pseudo-)randomness

```
def G(z):  
    if z > theta:  
        return "Heads"  
    else:  
        return "Tails"
```



Deterministic mapping



Sample from  
 $p(x)$

# Sampling

How do we generate a sample from a distribution?

Method 1: Go out and take a picture.  
(slow, expensive)



- $p(\mathbf{x})$  a distribution assigning some small probability to each possible setting of pixel values in a 24-bit-color 100x100 image.
- $\mathbf{x}$  a random variable, equal to some 100x100 image.

# Sampling

How do we generate a sample from a distribution?

Method 2: *Simulate* a sample.

```
np.random.uniform(0,1)
```



Source of (pseudo-)randomness

```
def G(z):  
    # hmmm...  
    # help?
```



Deterministic mapping



Sample from  
 $p(x)$

# How to write **G**?

It's just a  
function: learn it!

```
np.random.uniform(0,1)
```




Source of (pseudo-)randomness

```
def G(z):  
    return MLP(z, theta)
```



Deterministic mapping with  
randomly initialized parameters



Sample from  
 $p(x)$

# Learn G

But with what loss function?

```
np.random.uniform(0,1)
```




Source of (pseudo-)randomness

```
def G(z):  
    return MLP(z, theta)
```



Deterministic mapping with  
randomly initialized parameters




Sample from  
 $p(x)$


# Learn **G**

How do you teach **G** with  
no (in, out) examples?


I know my answer  
isn't right, but what  
should it have been  
for  $z=0.3$ ?



# Get two students (neural networks)



# Get two students (neural networks)






# Get two students (neural networks)

You're failing to fool the discriminator. Try making the top left pixels darker.


$z = 0.3$




# Get two students (neural networks)

You're failing to fool the discriminator. Try making the top left pixels darker.


$z = 0.3$



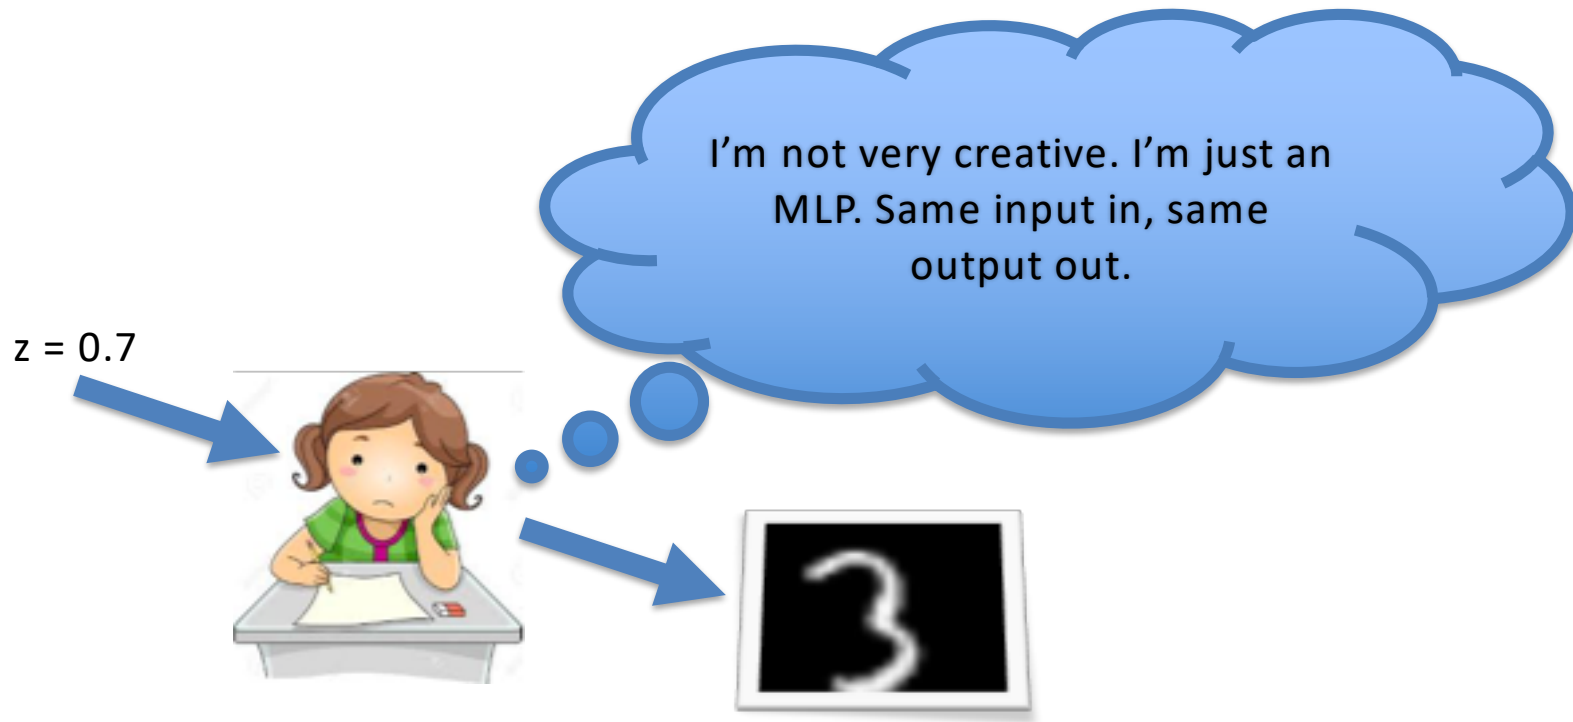
# Get two students (neural networks)



# Many iterations later...




# The importance of $z$



# The importance of $z$



# The importance of $z$



$G$  a *deterministic function*  
from  $z$  to images


# The importance of $z$




What if  $G$  learns to “play it safe” and always generate similar images?



# What if G always writes the same thing?




# What if G always writes the same thing?



# Main takeaway



We don't know how to teach G to draw. But with two students, we can make them teach each other!



# Implementation subtleties




```
z = tf.placeholder([50, 20])
```


- We provide *multiple* random numbers (say, 20) to inform G in creating *each* image.
- We ask G to create a whole “batch” (say, 50) of fake images each batch.

# Architecture of $G$ : $20 \rightarrow 256 \rightarrow 784$

$W_{G1} = \text{tf.Variable}(\text{tf.truncated\_normal}([20, 256], \text{stddev}=0.1))$



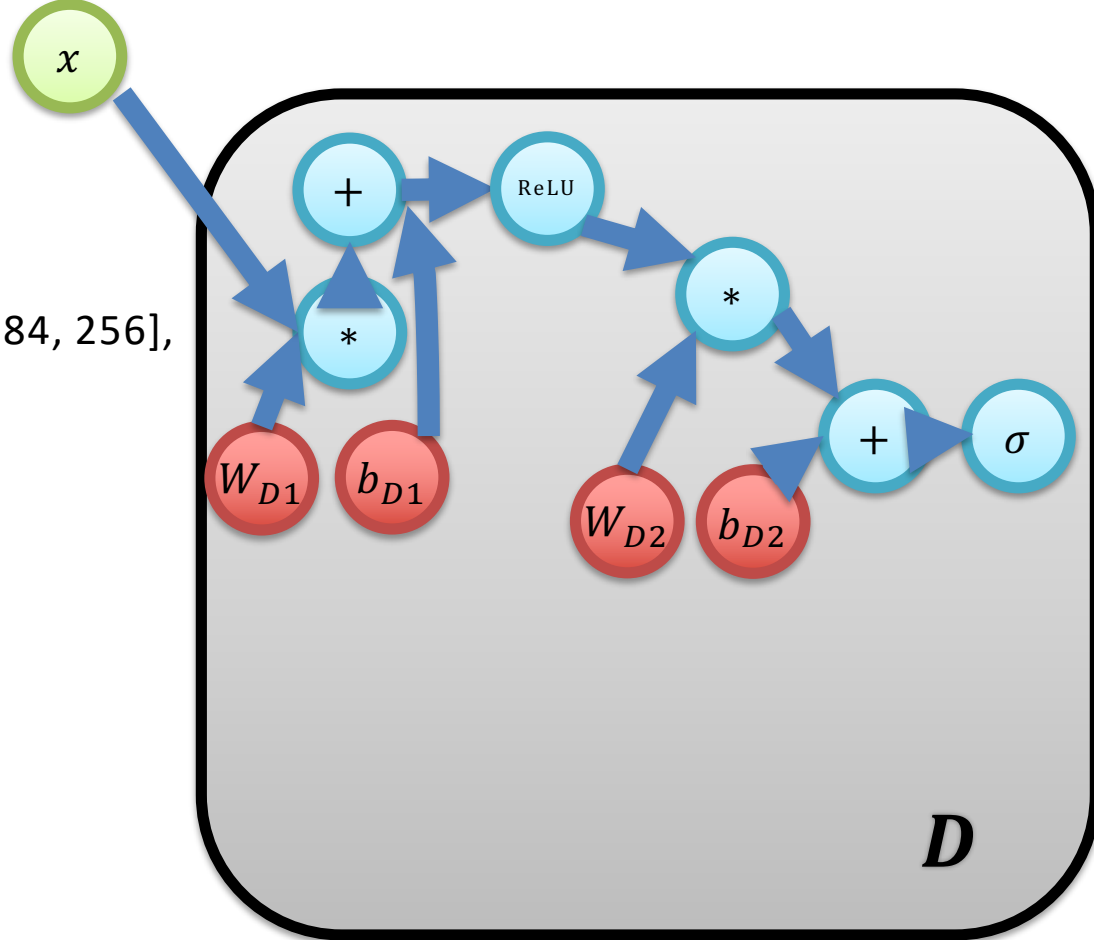
Our output: values between 0 and 1  
(black and white) for each pixel.

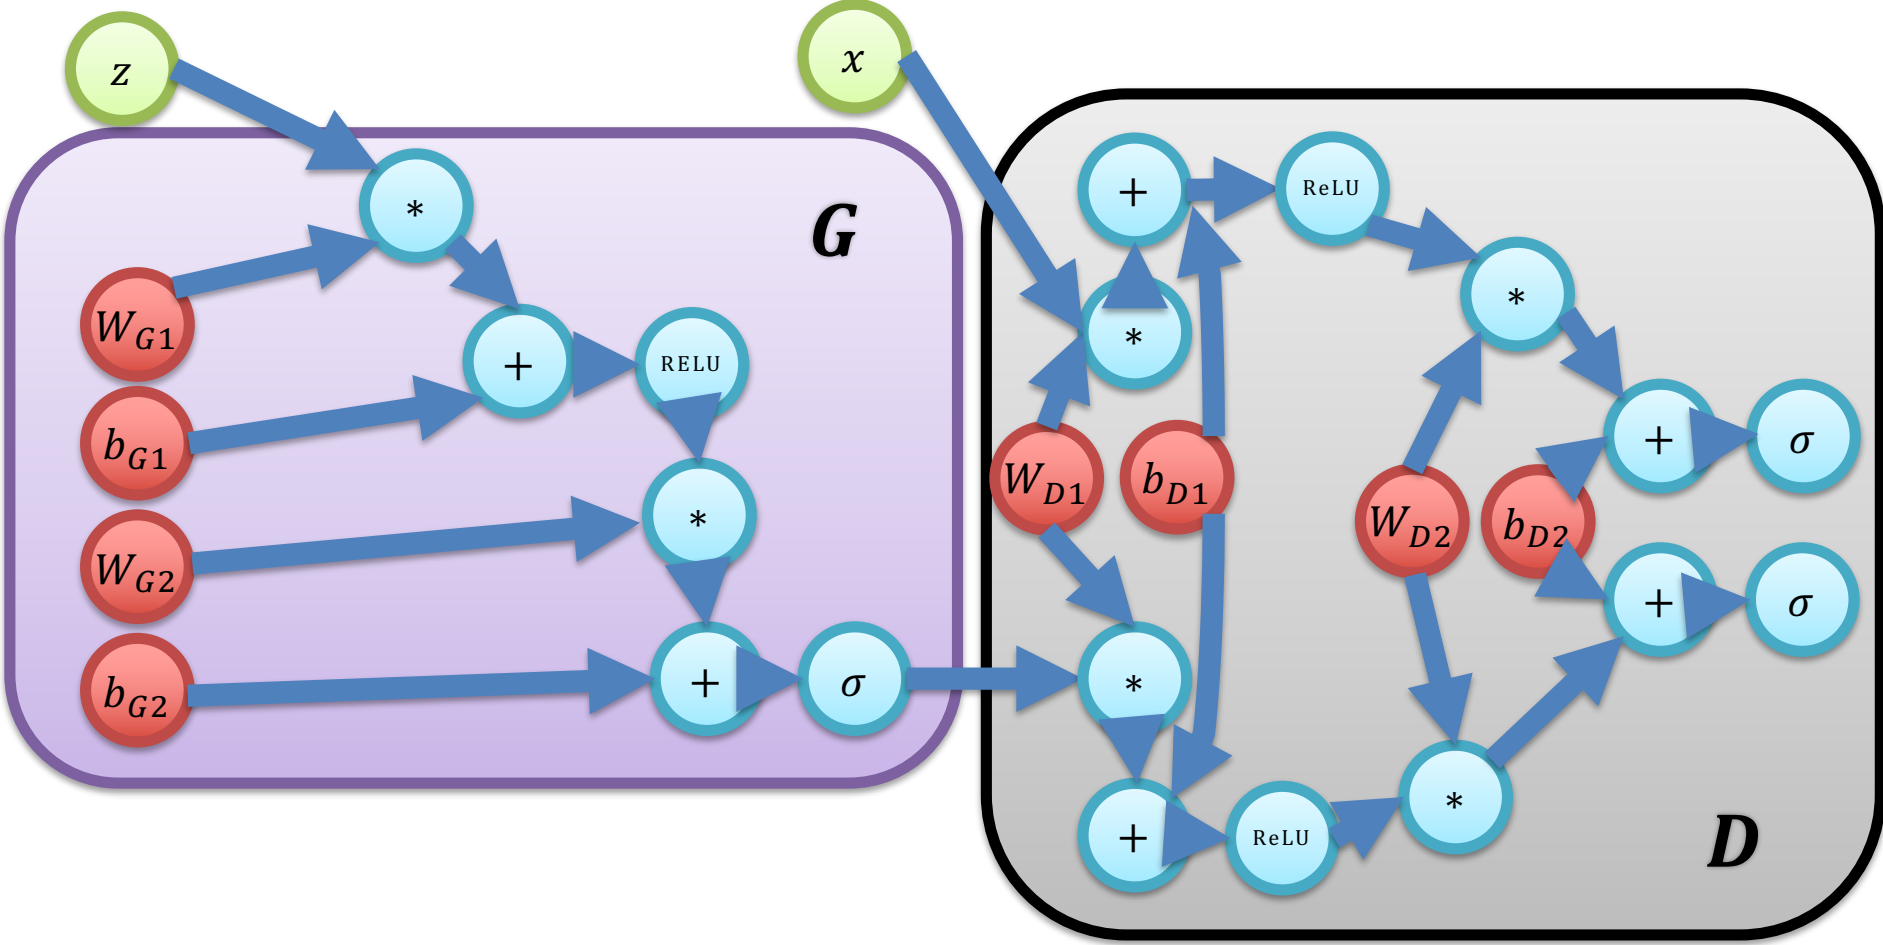


# Architecture of **D**:


784  $\rightarrow$  256  $\rightarrow$  1

$W_{D1} = \text{tf.Variable}(\text{tf.truncated\_normal}([784, 256],$   
 $\text{stddev}=0.1)$

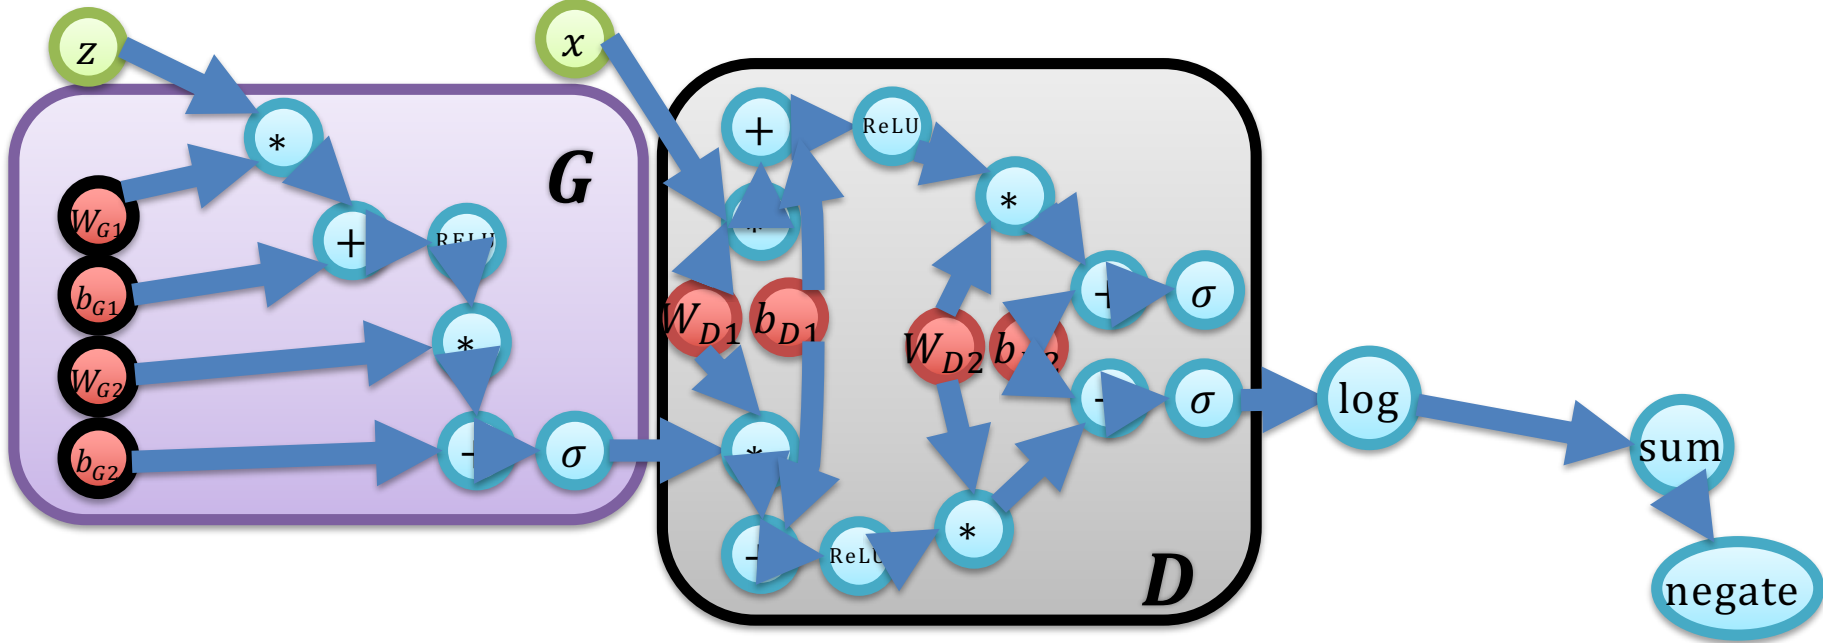









Loss when training  $D$



Loss when training  $G$



Loss when training  $G$